

DAY 4 - PostgreSQL 18 Backup & Disaster Recovery

LIVE Patroni / Zalando Postgres Operator Cluster - SINGLE-LINE COMMAND EDITION

Environment: namespace **postgres-ha** | cluster **acid-postgres-ha** | Docker Kubernetes | GCS bucket **postgres-ha-backups-parameswari-2026**

COPY/PASTE RULE: Every shell command below is one physical command line. Configuration files are shown separately as file contents and are not shell commands.

PostgreSQL 18 Backup & Disaster Recovery

Step-by-Step Lab Manual for a LIVE Patroni / Zalando Postgres Operator Cluster Namespace postgres-ha Cluster acid-postgres-ha (Zalando Operator + Patroni) PostgreSQL 18 Backup tools pg_dump, pg_dumpall, pg_basebackup, WAL archiving, pgBackRest Enterprise backup Full / Incremental / Differential / Compression / Encryption / Retention / PITR Restore rule Restore only to isolated recovery storage / separate recovery cluster CRITICAL SAFETY RULE: Take backups from the live cluster, but never overwrite the active Patroni member PVCs during restore or PITR. Perform physical restore only on isolated recovery storage or a separate recovery cluster. COPY/PASTE RULE: Every shell command in this corrected edition is a separate paragraph. Run commands one at a time. Do not paste two commands onto the same shell line unless they are intentionally joined with &&.

Table of Contents

- 0. Lab Architecture, Scope and Safety
- 1. Pre-Lab Validation of the Live Patroni Cluster
- 2. Backup Concepts: Logical, Physical, WAL and PITR
- 3. Lab 1 - Logical Backup from the Live Cluster
- 4. Lab 2 - Physical Backup with pg_basebackup (Corrected)
- 5. Lab 3 - WAL Archiving Validation
- 6. Lab 4 - Configure pgBackRest for the Live Cluster
- 7. Lab 5 - Full Backup with pgBackRest
- 8. Lab 6 - Incremental Backup
- 9. Lab 7 - Differential Backup
- 10. Repository, Compression, Encryption and Retention
- 11. Lab 8 - Restore / Full Cluster Recovery
- 12. Lab 9 - Point-in-Time Recovery (PITR)
- 13. Patroni Failover Considerations During Backup
- 14. Validation, Monitoring, Troubleshooting and Cleanup
- 15. RPO / RTO and Suggested Backup Schedule
- 16. Final Coverage Checklist

Implementation note: Zalando/Spilo commonly uses WAL-G/WAL-E for native backup. pgBackRest is not automatically present in all Spilo images. The pgBackRest portion therefore assumes an immutable image that contains a PostgreSQL 18-compatible pgBackRest build and a durable repository. Do not install packages ad hoc into a running production pod and assume they will survive restart.

0. Lab Architecture, Scope and Safety

0.1 What runs live and what is isolated?

- LIVE: acid-postgres-ha remains Patroni-managed and continues serving traffic during backup exercises.
- LIVE BACKUP: logical dumps, pg_basebackup, WAL archive checks and pgBackRest backup commands

run against the current primary or an eligible replica.

- ISOLATED RESTORE: physical restore and PITR use separate recovery storage. Never replace files below active PGDATA while Patroni is running.

- REPOSITORY: use durable object storage (GCS/S3) or, for training only, an RWX volume reachable by every member.

0.2 Recovery architecture

Application -> acid-postgres-ha Service -> Patroni Primary
Primary -> logical backup / physical backup / WAL archive / pgBackRest repository
Repository + archived WAL -> isolated recovery PGDATA -> validation -> controlled cutover

0.3 Why failover is not backup

Patroni protects availability. It does not protect against logical mistakes such as DROP TABLE, accidental DELETE, bad application updates, corruption that propagates to replicas, or retention mistakes. Backup and PITR provide recovery from those events.

1. Pre-Lab Validation of the Live Patroni Cluster

1.1 Set namespace and inspect resources

kubectl config set-context --current --namespace=postgres-ha
kubectl get pods -o wide
kubectl get svc
kubectl get postgresql
kubectl get pvc
kubectl get statefulset

1.2 Set variables

NS=postgres-ha
CLUSTER=acid-postgres-ha

1.3 Identify the current primary - do not assume pod-0

for p in acid-postgres-ha-0 acid-postgres-ha-1 acid-postgres-ha-2; do echo "=== \$p ==="; kubectl -n "\$NS" exec "\$p" -- curl -s http://127.0.0.1:8008/ jq -r '"role=" + (.role // "unknown") + " state=" + (.state // "unknown")'; done

Exactly one member should report primary/leader. If the cluster is unhealthy, stop the backup lab and repair HA first.

1.4 Set PRIMARY and REPLICA after checking

PRIMARY=acid-postgres-ha-0
REPLICA=acid-postgres-ha-1

Change both values to match the actual current primary and a healthy replica.

1.5 Validate role, version, PGDATA and LSN

kubectll -n "\$NS" exec "\$PRIMARY" -- psql -U postgres -Atqc "SELECT pg_is_in_recovery(), current_database(), version();"
kubectll -n "\$NS" exec "\$PRIMARY" -- psql -U postgres -Atqc "SHOW data_directory;"
kubectll -n "\$NS" exec "\$PRIMARY" -- psql -U postgres -c "SELECT now(), pg_current_wal_lsn(), pg_is_in_recovery();"
kubectll -n "\$NS" exec "\$PRIMARY" -- psql -U postgres -c "SELECT application_name, client_addr, state, sync_state, sent_lsn, write_lsn, flush_lsn, replay_lsn FROM pg_stat_replication;"

2. Backup Concepts: Logical, Physical, WAL and PITR

Logical backup (pg_dump/pg_dumpall) stores SQL objects and data and supports database/object-level restore. Physical backup (pg_basebackup/pgBackRest) copies the cluster data files and is the foundation for cluster-level recovery and PITR. WAL archiving continuously preserves transaction history. PITR restores a physical backup and replays archived WAL up to a selected time/LSN/transaction.

2.1 Full vs differential vs incremental

- Full: complete physical backup; establishes an independent backup baseline.
- Differential: changes since the most recent full backup. Restore uses that full + the chosen differential +

required WAL.

- Incremental: changes since the immediately preceding backup. Restore can require a chain of dependent

backups + required WAL.

- Archived WAL is not a substitute for a base backup; it is the continuous roll-forward stream used after

restoring a physical backup.

2.2 RPO and RTO

RPO is acceptable data loss. Frequent, healthy WAL archiving reduces potential data loss. RTO is acceptable service restoration time; it depends on backup size, storage/network speed, number of backup dependencies, WAL replay volume and validation time.

3. Lab 1 - Logical Backup from the Live Cluster

3.1 Create local backup directory

mkdir -p backups/logical

3.2 Full cluster logical backup

kubectll -n "\$NS" exec "\$PRIMARY" -- pg_dumpall -U postgres --clean --if-exists > "backups/logical/\${CLUSTER}-dumpall-\$(date +%Y%m%d-%H%M%S).sql"
ls -lh backups/logical/

3.3 Custom-format backup of banking database

```
kubectl -n "$NS" exec "$PRIMARY" -- pg_dump -U postgres -Fc -d banking -f /tmp/banking.dump

kubectl -n "$NS" cp "$PRIMARY:/tmp/banking.dump" "backups/logical/banking-$(date +%Y%m%d-%H%M%S).dump"
```

3.4 Verify without restoring

```
BACKUP_FILE=$(ls -lt backups/logical/banking-*.dump | head -1)

echo "$BACKUP_FILE"

pg_restore -l "$BACKUP_FILE" | head -40

SQL_FILE=$(ls -lt backups/logical/${CLUSTER}-dumpall-*.sql | head -1)

head -30 "$SQL_FILE"
```

3.5 Test logical restore into an isolated test database

```
kubectl -n "$NS" exec "$PRIMARY" -- psql -U postgres -c "DROP DATABASE IF EXISTS banking_restore_test;"

kubectl -n "$NS" exec "$PRIMARY" -- psql -U postgres -c "CREATE DATABASE banking_restore_test;"

BACKUP_FILE=$(ls -lt backups/logical/banking-*.dump | head -1)

echo "$BACKUP_FILE"

kubectl -n "$NS" cp "$BACKUP_FILE" "$PRIMARY:/tmp/banking-restore.dump"

kubectl -n "$NS" exec "$PRIMARY" -- pg_restore -U postgres -d banking_restore_test --clean --if-exists /tmp/banking-restore.dump

kubectl -n "$NS" exec "$PRIMARY" -- psql -U postgres -d banking_restore_test -c "\dt"
```

On a real production system, run restore tests in a dedicated recovery cluster instead of consuming I/O/storage on the production primary.

4. Lab 2 - Physical Backup with pg_basebackup (Corrected)

This section corrects the original manual. pg_basebackup uses the PostgreSQL replication protocol. A normal superuser login is not enough if pg_hba.conf does not allow that user for replication. In this Zalando/Patroni cluster, use the cluster primary Service and the operator-managed replication/standby credentials.

4.1 Verify the primary Service name

```
kubectl -n "$NS" get svc

PRIMARY_SERVICE=acid-postgres-ha

kubectl -n "$NS" exec "$REPLICA" -- getent hosts "$PRIMARY_SERVICE"
```

Do not use -h 127.0.0.1 when the client command is running in a replica pod. 127.0.0.1 means that same replica pod. Do not use -h postgres-ha unless that Service actually exists. For this cluster, the primary Service is acid-postgres-ha.

4.2 Inspect replication state and lag

```
kubectl -n "$NS" exec "$PRIMARY" -- psql -U postgres -c "SELECT application_name, client_addr, state, sync_state, sent_lsn, write_lsn, flush_lsn, replay_lsn, pg_size_pretty(pg_wal_lsn_diff(pg_current_wal_lsn(), replay_lsn)) AS replay_lag FROM pg_stat_replication;"
```

4.3 Identify replication-capable roles

```
kubectl -n "$NS" exec "$PRIMARY" -- psql -U postgres -c "SELECT rolname, rolreplication, rolcanlogin FROM pg_roles WHERE rolreplication = true;"
```

4.4 Check replication HBA rules

```
kubectl -n "$NS" exec "$PRIMARY" -- psql -U postgres -c "SELECT line_number, type, database, user_name, address, auth_method, error FROM pg_hba_file_rules WHERE 'replication' = ANY(database);"
```

Do not add broad rules such as host replication all 0.0.0.0/0 trust. Use the operator-managed replication role and existing HBA policy.

4.5 Find the operator-managed replication/standby secret

```
kubectl -n "$NS" get secrets | grep -Ei "standby|replic"
```

If the cluster exposes the standard Zalando standby credential secret, set it as follows. If the actual secret name differs, substitute the name returned by the previous command.

```
REPL_SECRET=standby.acid-postgres-ha.credentials.postgresql.acid.zalan.do
```

4.6 Load replication username and password into your local shell

```
REPL_USER=$(kubectl -n "$NS" get secret "$REPL_SECRET" -o jsonpath='{.data.username}' | base64 -d)
```

```
REPL_PASS=$(kubectl -n "$NS" get secret "$REPL_SECRET" -o jsonpath='{.data.password}' | base64 -d)
```

```
echo "$REPL_USER"
```

Do not echo the password during training or production. Keep it only in the shell variable and unset it after the exercise.

4.7 Test ordinary connectivity with the replication user

```
kubectl -n "$NS" exec "$REPLICA" -- env PGPASSWORD="$REPL_PASS" psql -h "$PRIMARY_SERVICE" -p 5432 -U "$REPL_USER" -d postgres -c "SELECT pg_is_in_recovery(), inet_server_addr(), inet_server_port();"
```

Expected: pg_is_in_recovery() = false because acid-postgres-ha routes to the current primary.

4.8 Prepare backup directory in the client pod

```
kubectl -n "$NS" exec "$REPLICA" -- rm -rf /tmp/basebackup
```

```
kubectl -n "$NS" exec "$REPLICA" -- mkdir -p /tmp/basebackup
```

4.9 Run pg_basebackup with the correct host and replication user

```
kubectl -n "$NS" exec "$REPLICA" -- env PGPASSWORD="$REPL_PASS" pg_basebackup -h "$PRIMARY_SERVICE" -p 5432 -U "$REPL_USER" -D /tmp/basebackup -Ft -z -Xs -P
```

Options: -Ft = tar format; -z = gzip; -Xs = stream WAL while the backup runs; -P = progress display.

4.10 Verify physical backup files

```
kubectl -n "$NS" exec "$REPLICA" -- ls -lh /tmp/basebackup
```

```
kubectl -n "$NS" exec "$REPLICA" -- find /tmp/basebackup -maxdepth 1 -type f -printf "%f %s bytes\n"
```

Typical tar-format output includes base.tar.gz and WAL tar output when WAL streaming is included.

4.11 Copy the backup out to the workstation

```
mkdir -p backups/physical

PHYSICAL_DIR="backups/physical/basebackup-$(date +%Y%m%d-%H%M%S)"

kubectl -n "$NS" cp "$REPLICA:/tmp/basebackup" "$PHYSICAL_DIR"

find "$PHYSICAL_DIR" -maxdepth 2 -type f -ls | head -20
```

4.12 Remove password variable from the shell

```
unset REPL_PASS
```

4.13 Troubleshooting the exact errors seen in this lab

• could not translate host name "postgres-ha": the Service name is wrong or does not exist. Use kubectl get svc and set PRIMARY_SERVICE=acid-postgres-ha.

• no pg_hba.conf entry for replication connection ... user "postgres": the host is reachable, but postgres is not permitted for replication. Use the operator-managed replication/standby role and its Secret.

• no pg_hba.conf entry ... SSL encryption / no encryption: the server evaluated SSL and non-SSL HBA paths and neither matched that replication user/source. Do not bypass this with trust; fix role/HBA/authentication correctly.

5. Lab 3 - WAL Archiving Validation

5.1 Inspect archive settings

```
kubectl -n "$NS" exec "$PRIMARY" -- psql -U postgres -c "SHOW archive_mode;"

kubectl -n "$NS" exec "$PRIMARY" -- psql -U postgres -c "SHOW archive_command;"

kubectl -n "$NS" exec "$PRIMARY" -- psql -U postgres -c "SHOW archive_library;"

kubectl -n "$NS" exec "$PRIMARY" -- psql -U postgres -x -c "SELECT * FROM pg_stat_archiver;"
```

5.2 Force a WAL switch and verify progress

```
BEFORE=$(kubectl -n "$NS" exec "$PRIMARY" -- psql -U postgres -Atqc "SELECT archived_count FROM pg_stat_archiver;")

kubectl -n "$NS" exec "$PRIMARY" -- psql -U postgres -c "SELECT pg_switch_wal();"

sleep 10

kubectl -n "$NS" exec "$PRIMARY" -- psql -U postgres -c "SELECT archived_count, failed_count, last_archived_wal, last_archive

echo "Archived before test: $BEFORE"
```

5.3 Check native Spilo backup tooling before changing archive_command

```
kubectl -n "$NS" exec "$PRIMARY" -- sh -lc 'command -v wal-g || true; command -v wal-e || true; env | grep -E "^(USE_WALG|WAL
```

Do not run two competing archive_command implementations without a migration plan. If WAL-G/WAL-E is active, decide whether pgBackRest will replace it or be used only for this controlled training path.

6. Lab 4 - Configure pgBackRest for Docker Kubernetes + GCS

Environment used in this corrected procedure: namespace postgres-ha, cluster acid-postgres-ha, pods acid-postgres-ha-0/1/2, Docker Kubernetes (not GKE), pgBackRest 2.59.1, GCS bucket postgres-ha-backups-parameswari-2026.

6.1 Record the current cluster configuration

NS=postgres-ha
CLUSTER=acid-postgres-ha
kubectl -n "\$NS" get pods -o wide
kubectl -n "\$NS" get postgresql "\$CLUSTER" -o yaml > "\${CLUSTER}-before-pgbackrest.yaml"
kubectl -n "\$NS" get statefulset "\$CLUSTER" -o yaml > "\${CLUSTER}-sts-before-pgbackrest.yaml"

6.2 Verify pgBackRest is already installed on every Patroni member

for pod in acid-postgres-ha-0 acid-postgres-ha-1 acid-postgres-ha-2; do; echo "==== \$pod ====="; kubectl -n postgres-ha exec "\$pod" -c postgres -- pgbackrest version; done

Expected: all three pods report pgBackRest 2.59.1. Because pgBackRest is already present in this environment, do not rebuild the Spilo image just to install pgBackRest. Lab 4 - pgBackRest for Docker Kubernetes + GCS (Corrected)

6.3 Verify the repository cipher Secret - do not recreate it if it already exists

The existing encrypted repository depends on the original cipher passphrase. If pgbackrest-repo-secret already exists, keep it. Do not generate a new passphrase just because a later validation fails.

kubectl -n postgres-ha get secret pgbackrest-repo-secret
--

Verify the expected key without printing its value. This direct-key JSONPath is supported by kubectl and replaces the earlier invalid range-variable example.

kubectl -n postgres-ha get secret pgbackrest-repo-secret -o jsonpath='{.data.repo-cipher-pass}' >/dev/null && echo 'repo-cipher-pass exists'
--

6.4 Docker Kubernetes: create/verify GCS service-account access

Docker Kubernetes cannot use repo1-gcs-key-type=auto because metadata.google.internal is not available. Use service-account key authentication for this training lab.

PROJECT_ID=\$(gcloud config get-value project)
SA_NAME=pgbackrest-gcs
SA_EMAIL="\${SA_NAME}@\${PROJECT_ID}.iam.gserviceaccount.com"
echo "\$PROJECT_ID"
echo "\$SA_EMAIL"
gcloud iam service-accounts describe "\$SA_EMAIL" >/dev/null 2>&1 gcloud iam service-accounts create "\$SA_NAME" --display-name="pgBackRest GCS backup"
gcloud storage buckets add-iam-policy-binding gs://postgres-ha-backups-parameswari-2026 --member="serviceAccount:\${SA_EMAIL}" --role="roles/storage.objectAdmin"

GCS key creation policy and Secret

6.4A If JSON key creation is blocked by Organization Policy If gcloud returns FAILED_PRECONDITION and constraints/iam.disableServiceAccountKeyCreation, the command is correct but the organization policy blocks new service-account keys.
GUI path: Google Cloud Console -> IAM & Admin -> Organization Policies -> search "Disable service account key creation". Open the constraint whose ID is iam.disableServiceAccountKeyCreation. Do not confuse it with

iam.disableServiceAccountKeyUpload or the managed constraint. At the project level, override the parent policy and set enforcement to Off / Not enforced. Wait for policy propagation.

```
gcloud org-policies describe constraints/iam.disableServiceAccountKeyCreation --project=postgresql-admin-lab-2026
```

Then retry key creation:

```
gcloud iam service-accounts keys create pgbackrest-gcs-key.json --iam-account="$SA_EMAIL"
```

```
chmod 600 pgbackrest-gcs-key.json
```

Create the Kubernetes Secret only if it does not already exist:

```
kubectl -n postgres-ha get secret pgbackrest-gcs-key || kubectl -n postgres-ha create secret generic pgbackrest-gcs-key --from-file=gcs-key.json=./pgbackrest-gcs-key.json
```

Verify only that the key exists; never print the JSON credential:

```
kubectl -n postgres-ha get secret pgbackrest-gcs-key -o jsonpath='{.data.gcs-key\.json}' >/dev/null && echo 'gcs-key.json exists'
```

Primary discovery and corrected pgBackRest configuration

6.5 Discover the current primary and PGDATA

Do not assume acid-postgres-ha-0 is primary. Determine the current primary from PostgreSQL itself.

```
PRIMARY=$(for pod in acid-postgres-ha-0 acid-postgres-ha-1 acid-postgres-ha-2; do role=$(kubectl -n postgres-ha exec "$pod" -c postgres -- psql -U postgres -Atqc "select case when pg_is_in_recovery() then 'replica' else 'primary' end;"); [ "$role" = "primary" ] && echo "$pod" && break; done)
```

```
echo "$PRIMARY"
```

```
PGDATA_ACTUAL=$(kubectl -n postgres-ha exec "$PRIMARY" -c postgres -- psql -U postgres -Atqc "SHOW data_directory;")
```

```
echo "$PGDATA_ACTUAL"
```

6.5A Create/update pgbackrest.conf for the fixed GCS bucket

File contents - pgbackrest.conf (paste into an editor; this block is not a shell command):

```
nano pgbackrest.conf
```

```
[postgres-ha]
pg1-path=/home/postgres/pgdata/pgroot/data
pg1-port=5432
[global]
repo1-type=gcs
repo1-gcs-bucket=postgres-ha-backups-parameswari-2026
repo1-gcs-key-type=service
repo1-gcs-key=/etc/pgbackrest-gcs/gcs-key.json
repo1-path=/pgbackrest/postgres-ha
repo1-retention-full=2
repo1-retention-diff=4
repo1-cipher-type=aes-256-cbc
compress-type=zst
compress-level=3
process-max=2
start-fast=y
log-level-console=info
```

```
[global:archive-push]
process-max=2
[global:archive-get]
process-max=2
```

If PGDATA_ACTUAL differs, replace pg1-path before updating the ConfigMap.

```
kubectl -n postgres-ha create configmap pgbackrest-config --from-file=pgbackrest.conf=./pgbackrest.conf --dry-run=client -o yaml | kubectl apply -f -
```

```
kubectl -n postgres-ha get configmap pgbackrest-config -o jsonpath='{.data.pgbackrest\.conf}' >/dev/null&& echo 'pgbackrest.conf exists'
```

Mount pgBackRest resources into every Patroni pod

6.6 Mount config, cipher Secret and GCS key

Use JSON rather than YAML for the merge patch. This avoids indentation loss in terminal paste dialogs.

File contents - pgbackrest-additional-volumes-patch.json (paste into an editor; this block is not a shell command):

```
nano pgbackrest-additional-volumes-patch.json

{
  "spec": {
    "additionalVolumes": [
      {
        "name": "pgbackrest-config",
        "mountPath": "/etc/pgbackrest",
        "targetContainers": ["postgres"],
        "volumeSource": {
          "configMap": {
            "name": "pgbackrest-config"
          }
        },
        "name": "pgbackrest-secret",
        "mountPath": "/etc/pgbackrest-secret",
        "targetContainers": ["postgres"],
        "volumeSource": {
          "secret": {
            "secretName": "pgbackrest-repo-secret"
          }
        },
        "name": "pgbackrest-gcs-key",
        "mountPath": "/etc/pgbackrest-gcs",
        "targetContainers": ["postgres"],
        "volumeSource": {
          "secret": {
            "secretName": "pgbackrest-gcs-key"
          }
        }
      ]
    }
  }
}
```

Dry run first. The file extension must be .json here; do not accidentally point kubectl at the old malformed .yaml file.

```
kubectl -n postgres-ha patch postgresql acid-postgres-ha --type merge --patch-file pgbackrest-additional-volumes-patch.json --dry-run=server -o yaml > /tmp/pgbackrest-patch-preview.yaml
```

```
kubectl -n postgres-ha patch postgresql acid-postgres-ha --type merge --patch-file pgbackrest-additional-volumes-patch.json
```

```
kubectl -n postgres-ha get pods -w
```

Terminal paste dialog: click Cancel if you need to edit. The warning dialog is not an editor. Prefer one-line commands when pasting.

Mount validation and cipher handling

6.7 Validate all mounts and files

```
kubectl -n postgres-ha get postgresql acid-postgres-ha -o jsonpath='{.spec.additionalVolumes}'; echo
```

```
kubectl -n postgres-ha get pod acid-postgres-ha-0 -o jsonpath='{range .spec.containers[?(@.name=="postgres")].volumeMounts[*]]{.name}{ " -> "}{.mountPath}{ "\n" }{end}'
```

Expected mounts: pgbackrest-config -> /etc/pgbackrest; pgbackrest-secret -> /etc/pgbackrest-secret; pgbackrest-gcs-key -> /etc/pgbackrest-gcs.

```
for pod in acid-postgres-ha-0 acid-postgres-ha-1 acid-postgres-ha-2; do; echo "==== $pod ====="; kubectl -n postgres-ha exec "$pod" -c postgres -- sh -c 'test -r /etc/pgbackrest/pgbackrest.conf && echo "config OK" || echo "config FAILED"; kubectl -n postgres-ha exec "$pod" -c postgres -- sh -c 'test -r /etc/pgbackrest-secret/repo-cipher-pass && echo "cipher OK" || echo "cipher FAILED"; kubectl -n postgres-ha exec "$pod" -c postgres -- sh -c 'test -r /etc/pgbackrest-gcs/gcs-key.json && echo "gcs-key OK" || echo "gcs-key FAILED"; done
```

6.8 Supply cipher password inside the same kubectl exec

The file exists inside the pod, not on the Ubuntu host. Never run cat /etc/pgbackrest-secret/... directly on the host.

```
kubectl -n postgres-ha exec "$PRIMARY" -c postgres -- sh -c 'export PGBACKREST_REPO1_CIPHER_PASS="$(cat /etc/pgbackrest-secret/repo-cipher-pass)"; pgbackrest --stanza=postgres-ha info'
```

If info reports "missing stanza path", GCS access is working and the next step is stanza-create.

Stanza creation - primary-aware and copy-safe

6.9 Create the stanza on the current PRIMARY only

First confirm roles. f = primary; t = replica.

```
for pod in acid-postgres-ha-0 acid-postgres-ha-1 acid-postgres-ha-2; do; echo "==== $pod ====="; kubectl -n postgres-ha exec "$pod" -c postgres -- psql -U postgres -Atqc "select pg_is_in_recovery();" ; done
```

Set PRIMARY to the pod that returns f, or rediscover it automatically:

```
PRIMARY=$(for pod in acid-postgres-ha-0 acid-postgres-ha-1 acid-postgres-ha-2; do v=$(kubectl -n postgres-ha exec "$pod" -c postgres -- psql -U postgres -Atqc "select pg_is_in_recovery();" ; [ "$v" = "f" ] &&echo "$pod" && break; done)
echo "$PRIMARY"
```

Run each pgBackRest action as a separate command. Do not paste stanza-create, check and info together without separators.

```
kubectl -n postgres-ha exec "$PRIMARY" -c postgres -- sh -c 'export PGBACKREST_REPO1_CIPHER_PASS="$(cat /etc/pgbackrest-secret/repo-cipher-pass)"; pgbackrest --stanza=postgres-ha stanza-create'
kubectl -n postgres-ha exec "$PRIMARY" -c postgres -- sh -c 'export PGBACKREST_REPO1_CIPHER_PASS="$(cat /etc/pgbackrest-secret/repo-cipher-pass)"; pgbackrest --stanza=postgres-ha check'
kubectl -n postgres-ha exec "$PRIMARY" -c postgres -- sh -c 'export PGBACKREST_REPO1_CIPHER_PASS="$(cat /etc/pgbackrest-secret/repo-cipher-pass)"; pgbackrest --stanza=postgres-ha info'
```

ERROR [056] unable to find primary cluster means stanza-create was executed against a replica. Re-detect PRIMARY and rerun only stanza-create.

Fix archive_command=/bin/true

6.10 Configure WAL archiving before Lab 5

If a full backup fails with ERROR [068] archive_command '/bin/true' must contain pgbackrest, pgBackRest is correctly installed but PostgreSQL WAL archiving is not configured for pgBackRest.

```
kubectl -n postgres-ha exec "$PRIMARY" -c postgres -- psql -U postgres -Atqc "show archive_mode;"
kubectl -n postgres-ha exec "$PRIMARY" -c postgres -- psql -U postgres -Atqc "show archive_command;"
```

Do not edit postgresql.conf inside a Patroni-managed pod. Merge the required settings through the Zalando PostgreSQL CR. First back up the CR:

```
kubectl -n postgres-ha get postgresql acid-postgres-ha -o yaml > acid-postgres-ha-before-archive-command.yaml
```

Create a merge patch that preserves existing PostgreSQL parameters and sets the pgBackRest archive command. The environment assignment is kept on the same command line so pgBackRest receives the cipher pass.

File contents - pgbackrest-archive-command-patch.json (paste into an editor; this block is not a shell command):

```
nano pgbackrest-archive-command-patch.json
{
  "spec": {
```

```
"postgresql": {
  "parameters": {
    "archive_mode": "on",
    "archive_command": "PGBACKREST_REPO1_CIPHER_PASS=$(cat /etc/pgbackrest-secret/repo-cipher-pass) pgbackrest --stanza=postgres-ha archive-push %p"
  }
}
}
}
}
```

```
kubectl -n postgres-ha patch postgresql acid-postgres-ha --type merge --patch-file pgbackrest-archive-command-patch.json --dry-run=server -o yaml > /tmp/archive-command-preview.yaml
```

```
kubectl -n postgres-ha patch postgresql acid-postgres-ha --type merge --patch-file pgbackrest-archive-command-patch.json
```

Wait for operator reconciliation, then rediscover PRIMARY because the leader can change during rolling restart.

Validate WAL archiving and run the full backup

6.11 Validate WAL archiving after reconciliation

```
kubectl -n postgres-ha get pods -w
```

```
PRIMARY=$(for pod in acid-postgres-ha-0 acid-postgres-ha-1 acid-postgres-ha-2; do v=$(kubectl -n postgres-ha exec "$pod" -c postgres -- psql -U postgres -Atqc "select pg_is_in_recovery();"); [ "$v" = "f" ] &&echo "$pod" && break; done)
```

```
echo "$PRIMARY"
```

```
kubectl -n postgres-ha exec "$PRIMARY" -c postgres -- psql -U postgres -Atqc "show archive_command;"
```

Expected: the command contains pgbackrest --stanza=postgres-ha archive-push %p and is no longer /bin/true.

```
kubectl -n postgres-ha exec "$PRIMARY" -c postgres -- psql -U postgres -c "select pg_switch_wal();"

sleep 10
```

```
kubectl -n postgres-ha exec "$PRIMARY" -c postgres -- psql -U postgres -x -c "select archived_count,failed_count,last_archived_wal,last_archived_time,last_failed_wal,last_failed_time from pg_stat_archiver;"
```

Lab 5 prerequisite: Lab 4 is complete only when stanza-create succeeds, pgbackrest check succeeds, pgbackrest info succeeds, archive_command contains pgbackrest, and pg_stat_archiver shows successful archiving. 7. Lab 5 - Full Backup with pgBackRest

```
kubectl -n postgres-ha exec "$PRIMARY" -c postgres -- sh -lc 'export PGBACKREST_REPO1_CIPHER_PASS="$(cat/etc/pgbackrest-secret/repo-cipher-pass)"; pgbackrest --stanza=postgres-ha --type=full backup'
```

```
kubectl -n postgres-ha exec "$PRIMARY" -c postgres -- sh -lc 'export PGBACKREST_REPO1_CIPHER_PASS="$(cat/etc/pgbackrest-secret/repo-cipher-pass)"; pgbackrest --stanza=postgres-ha info'
```

Expected: the newest backup label ends with F. Record backup start/stop time, database size, repository size and WAL start/stop values. Incremental and differential backups

8. Lab 6 - Incremental Backup

```
kubectl -n postgres-ha exec "$PRIMARY" -c postgres -- psql -U postgres -d banking -c "INSERT INTO dr_lab_events(note) VALUES ('before incremental backup');"

kubectl -n postgres-ha exec "$PRIMARY" -c postgres -- sh -lc 'export PGBACKREST_REPO1_CIPHER_PASS="$(cat/etc/pgbackrest-secret/repo-cipher-pass)"; pgbackrest --stanza=postgres-ha --type=incr backup; pgbackrest --stanza=postgres-ha info'
```

Expected: the new label ends with I and references an earlier backup set. Incremental = changes since the previous backup.

9. Lab 7 - Differential Backup

```
kubectl -n postgres-ha exec "$PRIMARY" -c postgres -- psql -U postgres -d banking -c "INSERT INTO dr_lab_events(note) VALUES ('before differential backup');"


```

```
kubectl -n postgres-ha exec "$PRIMARY" -c postgres -- sh -lc 'export PGBACKREST_REPO1_CIPHER_PASS="$(cat/etc/pgbackrest-secret/repo-cipher-pass)"; pgbackrest --stanza=postgres-ha --type=diff backup; pgbackrest --stanza=postgres-ha info'
```

Expected: the new label ends with D. Differential = changes since the most recent full backup.

10. Repository, Compression, Encryption and Retention

The repository stores backup sets, metadata and archived WAL. This lab uses GCS bucket postgres-ha-backups-parameswari-2026 so backups survive Kubernetes pod/node replacement.

```
compress-type=zst
compress-level=3
repo1-cipher-type=aes-256-cbc
repo1-retention-full=2
repo1-retention-diff=4
```

Keep the original repo cipher passphrase. Losing or replacing it makes encrypted backups unrecoverable.

10.3 Encryption

```
repo1-cipher-type=aes-256-cbc
```

The cipher passphrase is supplied from the Kubernetes Secret via PGBACKREST_REPO1_CIPHER_PASS. Losing the passphrase makes encrypted backups unrecoverable.

10.4 Retention

```
repo1-retention-full=2
repo1-retention-diff=4
```

10.5 Review expiration before deleting backups

```
kubectl -n "$NS" exec "$PRIMARY" -- sh -lc 'export PGBACKREST_REPO1_CIPHER_PASS="$(cat /etc/pgbackrest-secret/repo-cipher-pass)"; pgbackrest --stanza=postgres-ha --dry-run expire'
kubectl -n "$NS" exec "$PRIMARY" -- sh -lc 'export PGBACKREST_REPO1_CIPHER_PASS="$(cat /etc/pgbackrest-secret/repo-cipher-pass)"; pgbackrest --stanza=postgres-ha expire'
```

11. Lab 8 - Restore / Full Cluster Recovery

DESTRUCTIVE AT PGDATA LEVEL: Never run this against active acid-postgres-ha PVCs. Create a separate recovery cluster/PVC first.

11.1 Create a separate recovery cluster

File contents - acid-pgbackrest-recovery.yaml (paste into an editor; this block is not a shell command):

```
nano acid-pgbackrest-recovery.yaml

apiVersion: acid.zalan.do/v1
kind: postgresql
metadata:
  name: acid-pgbackrest-recovery
  namespace: postgres-ha
spec:
  teamId: acid
  numberOfInstances: 1
  volume:
```

```
size: 20Gi
postgresql:
  version: "18"
  dockerImage: registry.example.com/spilo-pgbackrest:pg18

kubectl apply -f acid-pgbackrest-recovery.yaml
```

11.2 Discover recovery StatefulSet/PVC and service account

```
REC_STS=acid-pgbackrest-recovery

kubectl -n "$NS" scale statefulset/"$REC_STS" --replicas=0

kubectl -n "$NS" get pvc | grep "$REC_STS"

REC_PVC=$(kubectl -n "$NS" get pvc -o name | grep acid-pgbackrest-recovery | head -1 | cut -d/ -f2)

echo "$REC_PVC"

POD_SA=$(kubectl -n "$NS" get pod "$PRIMARY" -o jsonpath='{.spec.serviceAccountName}')

echo "$POD_SA"
```

11.3 Create recovery toolbox pod

File contents - pgbackrest-recovery-toolbox.yaml (paste into an editor; this block is not a shell command):

```
nano pgbackrest-recovery-toolbox.yaml

apiVersion: v1
kind: Pod
metadata:
  name: pgbackrest-recovery-toolbox
  namespace: $NS
spec:
  serviceAccountName: $POD_SA
  restartPolicy: Never
  containers:
  - name: toolbox
    image: registry.example.com/spilo-pgbackrest:pg18
    command: ["sh","-lc","sleep infinity"]
    volumeMounts:
    - name: data
      mountPath: /home/postgres/pgdata
    - name: config
      mountPath: /etc/pgbackrest
    - name: secret
      mountPath: /etc/pgbackrest-secret
    volumes:
    - name: data
      persistentVolumeClaim:
        claimName: $REC_PVC

  - name: configconfigMap:name: pgbackrest-config- name: secretsecret:secretName: pgbackrest-repo-secret EOF
```

```
kubectl apply -f pgbackrest-recovery-toolbox.yaml

kubectl -n "$NS" wait --for=condition=Ready pod/pgbackrest-recovery-toolbox --timeout=10m
```

11.4 Discover and prepare recovery PGDATA

```
kubectl -n "$NS" exec pgbackrest-recovery-toolbox -- sh -lc 'find /home/postgres/pgdata -maxdepth 4 -name PG_VERSION -print'

REC_PGDATA=/home/postgres/pgdata/pgroot/data
```

Verify REC_PGDATA before running the next command.

```
kubectl -n "$NS" exec pgbackrest-recovery-toolbox -- sh -lc "mv "$REC_PGDATA" "${REC_PGDATA}.before-restore" && mkdir -p "$REC_PGDATA"
```

11.5 Restore the latest backup

```
kubectl -n "$NS" exec pgbackrest-recovery-toolbox -- sh -lc 'export PGBACKREST_REPO1_CIPHER_PASS="$(cat /etc/pgbackrest-secret/repo-cipher-pass)"; pgbackrest --stanza=postgres-ha --pg1-path=/home/postgres/pgdata/pgroot/data restore'
```

Correction from the previous manual: do not use --set=latest. To restore the latest eligible backup, omit --set. Use --set only when you have an actual pgBackRest backup label.

11.6 Start and validate the recovery cluster

```
kubectl -n "$NS" delete pod pgbackrest-recovery-toolbox --wait=true

kubectl -n "$NS" scale statefulset/"$REC_STS" --replicas=1

kubectl -n "$NS" rollout status statefulset/"$REC_STS" --timeout=10m

REC_POD=$(kubectl -n "$NS" get pod -l cluster-name=acid-pgbackrest-recovery -o jsonpath='{.items[0].metadata.name}')

kubectl -n "$NS" exec "$REC_POD" -- psql -U postgres -Atqc "SELECT pg_is_in_recovery(), now();"

kubectl -n "$NS" exec "$REC_POD" -- psql -U postgres -d banking -c "SELECT * FROM dr_lab_events ORDER BY id;"
```

12. Lab 9 - Point-in-Time Recovery (PITR)

12.1 Take a fresh full backup

```
kubectl -n "$NS" exec "$PRIMARY" -- sh -lc 'export PGBACKREST_REPO1_CIPHER_PASS="$(cat /etc/pgbackrest-secret/repo-cipher-pass)"; pgbackrest --stanza=postgres-ha --type=full backup;

pgbackrest --stanza=postgres-ha info'
```

12.2 Create a known-good recovery point

```
kubectl -n "$NS" exec "$PRIMARY" -- psql -U postgres -d banking -c "INSERT INTO dr_lab_events(note) VALUES ('PITR good row');"

TARGET_TIME=$(kubectl -n "$NS" exec "$PRIMARY" -- psql -U postgres -Atqc "SELECT to_char(clock_timestamp() AT TIME ZONE 'UTC', 'YYYY-MM-DD HH24:MI:SS.US') || '+00';")

echo "$TARGET_TIME"

sleep 3
```

12.3 Simulate accidental data loss

```
kubectl -n "$NS" exec "$PRIMARY" -- psql -U postgres -d banking -c "DELETE FROM dr_lab_events;"

kubectl -n "$NS" exec "$PRIMARY" -- psql -U postgres -d banking -c "SELECT count(*) FROM dr_lab_events;"

kubectl -n "$NS" exec "$PRIMARY" -- psql -U postgres -c "SELECT pg_switch_wal();"
```

12.4 Prepare isolated recovery target again

```
kubectl -n "$NS" scale statefulset/"$REC_STS" --replicas=0

kubectl -n "$NS" wait --for=delete pod/"$REC_POD" --timeout=180s || true

kubectl apply -f pgbackrest-recovery-toolbox.yaml

kubectl -n "$NS" wait --for=condition=Ready pod/pgbackrest-recovery-toolbox --timeout=10m
```

12.5 Restore to target time

Pass TARGET_TIME into the recovery pod with env so the local value is available to the remote shell without unsafe quoting or local command substitution.

```
kubectl -n "$NS" exec pgbackrest-recovery-toolbox -- env TARGET_TIME="$TARGET_TIME" sh -lc 'export

PGBACKREST_REPO1_CIPHER_PASS="$(cat /etc/pgbackrest-secret/repo-cipher-pass)"; pgbackrest --stanza=postgres-ha --pg1-path=/home/postgres/pgdata/pgroot/data --type=time --target="$TARGET_TIME" --target-action=promote restore'
```

12.6 Start and validate PITR result

```
kubectl -n "$NS" delete pod pgbackrest-recovery-toolbox --wait=true

kubectl -n "$NS" scale statefulset/"$REC_STS" --replicas=1

kubectl -n "$NS" rollout status statefulset/"$REC_STS" --timeout=10m

REC_POD=$(kubectl -n "$NS" get pod -l cluster-name=acid-pgbackrest-recovery -o
jsonpath='{.items[0].metadata.name}')

kubectl -n "$NS" exec "$REC_POD" -- psql -U postgres -d banking -c "SELECT * FROM dr_lab_events ORDER BY id;"

kubectl -n "$NS" exec "$REC_POD" -- psql -U postgres -c "SELECT pg_is_in_recovery();"

kubectl -n "$NS" exec "$REC_POD" -- psql -U postgres -c "SELECT timeline_id FROM pg_control_checkpoint();"
```

Expected: rows committed before TARGET_TIME are present, the accidental DELETE is absent, and pg_is_in_recovery() is false after promotion.

13. Patroni Failover Considerations During Backup

Rediscover the current primary before each critical backup or archive action. A stale PRIMARY variable can point to a former primary after switchover/failover.

```
PRIMARY=$(for p in $(kubectl -n "$NS" get pod -l cluster-name="$CLUSTER" -o name | cut -d/ -f2); do role=$(kubectl -n "$NS" exec "$p" -- curl -s http://127.0.0.1:8008/ | jq -r '.role'); if [ "$role" = "primary" ] || [ "$role" = "master" ] || [ "$role" = "leader" ]; then echo "$p"; fi; done | head -1)

echo "Current primary: $PRIMARY"

kubectl -n "$NS" exec "$PRIMARY" -- psql -U postgres -c "SELECT pg_switch_wal();"

kubectl -n "$NS" exec "$PRIMARY" -- sh -lc 'export PGBACKREST_REPO1_CIPHER_PASS="$(cat /etc/pgbackrest-secret/repo-cipher-pass)"; pgbackrest --stanza=postgres-ha check; pgbackrest --stanza=postgres-ha info'
```

14. Validation, Monitoring, Troubleshooting and Cleanup

14.1 Daily validation

kubectll -n "\$NS" exec "\$PRIMARY" -- psql -U postgres -x -c "SELECT * FROM pg_stat_archiver;"
kubectll -n "\$NS" exec "\$PRIMARY" -- sh -lc 'export PGBACKREST_REPO1_CIPHER_PASS="\$(cat /etc/pgbackrest-secret/repo-cipher-pass)"; pgbackrest --stanza=postgres-ha check; pgbackrest --stanza=postgres-ha info'
kubectll -n "\$NS" exec "\$PRIMARY" -- patronictl list

14.2 Cleanup only recovery objects

kubectll -n "\$NS" delete postgresql acid-pgbackrest-recovery --ignore-not-found
kubectll -n "\$NS" delete pod pgbackrest-recovery-toolbox --ignore-not-found
kubectll -n "\$NS" get pvc grep acid-pgbackrest-recovery true

Do not delete acid-postgres-ha PVCs or the pgBackRest repository as part of generic cleanup.

15. RPO / RTO and Suggested Backup Schedule

- WAL archive: continuous; determines how close to the incident you can recover.
- Full: weekly baseline/anchor, adjusted to database size and RTO.
- Differential: daily if you want fewer restore dependencies than a long incremental chain.
- Incremental: every 4-6 hours or as workload/RPO requires.
- Logical pg_dump: daily/weekly for critical databases when object-level recovery is useful.
- Restore drill: at least monthly and after major PostgreSQL/Spilo/pgBackRest/storage/IAM/operator

changes.

15.1 Backup commands

pgbackrest --stanza=postgres-ha --type=full backup
pgbackrest --stanza=postgres-ha --type=diff backup
pgbackrest --stanza=postgres-ha --type=incr backup

16. Final Coverage Checklist

Topic Status Logical backup Covered Physical backup with pg_basebackup Covered - corrected host/user/HBA/secret flow WAL archiving Covered PITR Covered pgBackRest configuration Covered Repository Covered Compression Covered Encryption Covered Incremental backup Covered Differential backup Covered Retention policies Covered Full backup Covered Restore Covered on isolated recovery storage Patroni failover handling Covered Copy/paste-friendly commands Every command is a separate paragraph Most important corrections in this edition: pg_basebackup now uses acid-postgres-ha plus the replication/standby Secret instead of 127.0.0.1 + postgres; the latest pgBackRest restore omits the invalid --set=latest pattern; recovery toolbox inherits the database pod service account for object-storage access; and every shell command is isolated for reliable copy/paste.